

PostgreSQL and Google Go Language


April 14, 2015, Austin, Texas, USA

John Scott

Consultant, American Messaging, 2006-now

Founder, SetSpace, Inc, 1998-now

Glenda the Gopher



About the Talk

- History, Motivation and Highlights of Go Language
- How to Query PostgreSQL in Go
- Benchmarks Comparing Go Query to C Code

About Go

Go is an open source programming language that makes it easy to build simple, reliable, and efficient software.

Design began in late 2007.

- Robert Griesemer, Rob Pike, Ken Thompson
- Russ Cox, Ian Lance Taylor

Became open source in November 2009.

Developed entirely in the open; very active community.

Language stable as of Go 1, early 2012.

Work continues.

Go

A deliberately simple but powerful and fun language.

- start with C, remove complex parts
- add interfaces, concurrency
- also: garbage collection, closures, reflection, slices ...
- strings encoded in UTF-8 only

For more background on design:

- [Less is exponentially more](http://commandcenter.blogspot.com/2012/06/less-is-exponentially-more.html) (<http://commandcenter.blogspot.com/2012/06/less-is-exponentially-more.html>)
- [Go at Google: Language Design in the Service of Software Engineering](http://talks.golang.org/2012/splash.article) (<http://talks.golang.org/2012/splash.article>)

Building Go Executables

- Statically compiled and linked executables
- No Runtime Libraries - Not Even libc!
- Size of Smallest Executable is Small (400KB)
- Process Startup about 50-100x Quicker Java
- Range Checking on Arrays/Slice
- Null Pointer Deferences Caught at Runtime
- No Magic

Production Go

- Docker blog.golang.org/docker (<http://blog.golang.org/docker>)
- Dropbox github.com/dropbox/godropbox (<https://github.com/dropbox/godropbox>)
- YouTube, Twitch
- GitHub, Tumblr (gocircuit!), Heroku
- Canonical/Ubuntu, Mozilla, SoundCloud
- Digital Ocean, Getty Images, DataDog
- Apple, Walmart, Intel are Recently Advertising Go Jobs
- Honest Dollar in Austin

For Jobs See [golangprojects.com](http://www.golangprojects.com/) (<http://www.golangprojects.com/>)

Go Syntax Easy to Carry in Your Head

- for {} is only loop construct
- if/else, switch{}, goto, labeled break
- common unary/binary operators: */+- == != > >= < <= ++ -- << >>
- address/content prefix operators: & * (with garbage collection!)
- json style struct/interface {} declarations
- arrays and slices (python)
- functions with closures (javascript)
- interface just a set of methods
- channels multiplex with select{} statement

Unified Declaration and Assignment

```
// create variable initialized to utf-8 string
// no need for 'var' statement

s := "hello, world"

i := 0

// explicit creation of 8 bit unsigned variable
ui := uint8(0)
```

New Types - Unified Declaration and Assignment

```
type bean_count uint16

// cast 0 to bean_count and create a variable
bc := bean_count(0)

Println(bc)
Println(uint16(bc))
Println(bean_count(0))
```

- Cast must be explicit but compiler is smart about type lineage

Methods - Unified Declaration and Assignment

```
type bean_count uint16

// bind private print() to any bean_count type

func (bc bean_count) print() {
    Println(bc)
}

...

bc := bean_count(0)
bc.print()
```

- A function bound to a type is called a "Method"

Structures - Unified Declaration and Assignment

```
type wine struct {  
    vitner    string  
    name      string  
    year      uint16  
}  
  
//  json style initialization of new struct value  
  
w := wine{  
    vitner:    "Conchay Toro",  
    name:      "Casillero del Diablo",  
    year:      2008,  
}  
w.year = 2013  
  
//  take address of wine structure value.  '&' operator from C language  
  
wp := &w  
  
//  dereference pointer using '.' and not the C style '->' operator  
  
wp.year = 2012
```

Function Pointers - Unified Declaration and Assignment

```
// variable info is a pointer to a function - named closure

info := func(msg string) {
    Println(msg)
}

info("hello, world")

...

info("good bye, cruel world")
```

Memory Reference - Unified Declaration and Assignment

```
// create variable initialized to utf-8 string  
  
s := "hello, world"  
  
// take memory address of variable s  
  
sp := &s  
  
// change contents of what new variable s references  
  
*sp = "good bye, cruel world"
```

Run Memory Reference - Unified Declaration and Assignment

```
package main

import "fmt"

func main() {

    // create variable initialized to utf-8 string

    s := "hello, world"
    fmt.Println(s)

    // take memory address of variable s

    sp := &s

    // change what s references

    *sp = "good bye, cruel world"

    fmt.Println(s)      // prove it
}
```

[Run](#)

Arrays and Slices - Unified Declaration and Assignment

- Arrays are static with immutable sizes
- Slices (from python) are windows into a static array
- Slices reference from 0
- Slices can grow and shrink but not beyond the capacity of array

Channels - Unified Declaration and Assignment

- Think of full duplex pipes for any data type ... including channels

```
done := make(chan int)

...

// send the boolean value 'true' down the channel.
// block until other ends reads

done <- true

...

// the other end waits for done

<- done
Print("finished")
```

Go is a Blend of Sequential and Concurrent Coding

Typical Sequential Coding

Fibonacci series: $f(n) = f(n-1) + f(n-2)$

```
package main
import . "fmt"

func fib(i int) int {

    if i < 2 {
        return 1
    }

    return fib(i - 1) + fib(i - 2)
}

func main() {

    Println(fib(40))

}
```

[Run](#)

Slower Sequential Coding

Fibonacci series: $f(n) = f(n-1) + f(n-2)$

```
package main
import . "fmt"

func fib(i int) int {

    if i < 2 {
        return 1
    }

    return fib(i - 1) + fib(i - 2)
}

func main() {

    Println(fib(40), fib(40), fib(40), fib(40))
}
```

[Run](#)

Easy Concurrent Coding

```
package main
import . "fmt"

func fib(i int) int {
    if i < 2 {
        return 1
    }
    return fib(i - 1) + fib(i - 2)
}

func main() {

    answer := make(chan int)

    fib_worker := func(i int) {
        answer <- fib(i)
    }

    go fib_worker(40)          // Just add "go" in front of function call
    go fib_worker(40)
    go fib_worker(40)
    go fib_worker(40)

    println(<- answer, <- answer, <- answer, <- answer)
}
```

[Run](#)

Concurrency with Channels

Share memory by communicating rather than communicating by sharing memory

- Based on Communicating Sequential Processes by Sir Tony Hoare
- Channels are Typed, Two Way Pipes
- Writer Blocks until Reader Completes
- Read on Closed Channel Returns Nil - A Cheap Broadcast
- Write on Closed Channel is Runtime Panic
- Multiple Channels Read/Written with `select{} statement`
- Channels can be Buffered

Any Datatype Sent Over a Channel

Channels can be sent over channels

```
type session struct {  
    name    string  
}  
  
c := make(chan *session)  
  
// channel of channels of pointers to sessions  
load_balancer := make(chan chan *session)
```

"Hello, World" is Load Balancer

Select Reads/Write Many Channels at Once

```
func flow (north, south, east, west chan int) {  
  
    var message int  
  
    for {  
        select {  
  
            case message <- west:  
                ..  
            case message <- south:  
                ..  
            case north <- message:  
                ..  
            case east <- message:  
                ..  
        }  
    }  
}
```


Go Memory Management

- Garbage Collection of Unreferenced Data
- True Pointers and Writing Directly into Memory
- Programmer Controls Memory Layout
- Possible to Write a Typed malloc(), Relieving Pressure on Collector

Database/Sql Core Package and PostgreSQL pq

GoLang Package database/sql

Abstracts SQL Databases

- Similar to famous Perl package named DBI
- Written by Brad Fitzpatrick (Live Journal)
- Two PostgreSQL Drivers: native or static link against C library libpq
- Other Drivers: MS Server, MySQL, SQLite, DB2, ODBC, Oracle8, Sybase, Firebird, YQL

Builtin Connection Pooling

- All databases drivers can pool
- Single DB Object transparently pools across connections
- Transaction binds to single pool connection
- SetMin/MaxIdleConnections()

database/sql Data Types

- Database
- Driver (referenced via URI)
- Statement
- Transaction
- Row, Rows,
- Result
- Bool, Float64, Int64, String

database/sql Setup Methods

- `Open()`
- `Prepare()`
- `Close()`
- `Begin()`
- `Commit()`
- `Rollback()`

Query Methods - Package database/sql

- Exec()
- Query()
- QueryRow()
- Next()
- Scan()

Example - Query Multiple Rows

```
rows, err := db.Query("SELECT ...")

...

defer rows.Close()

for rows.Next() {
    var id int
    var name string

    err = rows.Scan(&id, &name)

    ...
}

err = rows.Err() // get any error encountered during iteration
...
```


PostgreSQL Driver database/sql/pq

- 100% go language - no linking of C library libpq
- Hidden driver underneath generic core library
- Not distributed with golang core - available on github
- Very, very fast

```
$ go get github.com/lib/pq
```

sql-bench.go - import statements

```
package main

import (
    "bufio"
    "database/sql"
    "io"
    "os"
    "strings"
    "syscall"

    . "fmt"
    _ "github.com/lib/pq"
)
```

sql-bench.go - constants & declarations

```
const (  
    select_stmt = `  
        select  
            exists (  
                select  
                    cookie  
                from  
                    login_session  
                where  
                    cookie = $1  
            )`  
    QUERY_COUNT = 9  
    COOKIE_CHANNEL_SIZE = 96  
)  
  
var (  
    Stdin = os.NewFile(uintptr(syscall.Stdin), "/dev/stdin")  
)  
  
type stat struct {           // track stats for select queries for each go routine  
    true_count    uint64  
    false_count   uint64  
}
```

sql-bench.go - open database and prepare statement

```
func open_db() (  
    db *sql.DB,  
    stmt *sql.Stmt,  
) {  
    var err error  
  
    // open postgresql database, inheriting PG params from environment  
    db, err = sql.Open("postgres", "sslmode=disable")  
    if err != nil {  
        panic(err)  
    }  
  
    db.SetMaxIdleConns(8)          // keep 8 idle connections  
    db.SetMaxOpenConns(128)       // no more 128 simultaneous db connection  
  
    stmt, err = db.Prepare(select_stmt)  
    if err != nil {  
        panic(err)  
    }  
    return db, stmt  
}
```

sql-bench.go - query worker

```
func select_cookies(  
    db *sql.DB,                // database connection  
    stmt *sql.Stmt,            // prepared statment  
    cookies chan string,       // read cookies  
    done chan *stat,           // answer with stats  
) {  
    stat, exists := &stat{}, false  
  
    // read cookies on string channel  
    for cookie := range cookies {  
        err := stmt.QueryRow(cookie).Scan(&exists)  
        if err != nil {  
            panic(err)  
        }  
        if exists {  
            stat.true_count++  
        } else {  
            stat.false_count++  
        }  
    }  
    done <- stat  
}
```

sql-bench.go - send cookies to competing workers

```
func send_cookies (cookies chan string) {  
  
    //  for each cookie read on standard input,  
    //  send that cookie to some query go routines  
    in := bufio.NewReader(Stdin)  
    for {  
        cookie, err := in.ReadString('\n')  
        if err != nil {  
            if err == io.EOF {  
                break  
            }  
            panic(err)  
        }  
  
        //  send to competing pool of selects  
        cookies <- strings.TrimRight(cookie, "\n")  
    }  
    close(cookies)  
}
```

sql-bench.go - wait for queries to finish, print stats

```
func wait(done chan *stat) {  
  
    true_count := uint64(0)  
    false_count := uint64(0)  
  
    // wait for select()s to finish in the  
    for i := 0; i < QUERY_COUNT; i++ {  
  
        // read stat from each terminating query worker  
        s := <- done  
  
        // accumulate total stats as workers exit  
        true_count += s.true_count  
        false_count += s.false_count  
    }  
    Printf(" True/False Count: %d/%d\n", true_count, false_count)  
}
```

sql-bench.go - run the benchmark in main()

```
func main() {  
  
    // open postgresql database, inheriting PG params from environment  
    db, stmt := open_db()  
  
    // works indicate finished by sending stats on 'done' channel  
    done := make(chan *stat)  
  
    // workers compete by reading string cookies from this channel  
    cookies := make(chan string, COOKIE_CHANNEL_SIZE)  
  
    // spawn competing go routines  
    for i := 0; i < QUERY_COUNT; i++ {  
        go select_cookies(db, stmt, cookies, done)  
    }  
    // boot up cookie switch in background  
    go send_cookies(cookies)  
  
    // synchronously wait for results  
    wait(done)  
  
    os.Exit(0)    // good bye, cruel world  
}
```


Benchmark Comparing Queries in Go/pq to C/libpq

Source Code for C and Go Benchmarks

Go Benchmark in `sql-bench.go`

<http:pgdfw-gopq/sql-bench.go> (<http:pgdfw-gopq/sql-bench.go>)

Shell Script to Run `sql-bench-go.sh`

<http:pgdfw-gopq/sql-bench-go-sh.txt> (<http:pgdfw-gopq/sql-bench-go-sh.txt>)

C Benchmark in `sql-bench.pgc`

<http:pgdfw-gopq/sql-bench-pgc.txt> (<http:pgdfw-gopq/sql-bench-pgc.txt>)

Shell Script to Run `sql-bench.pgc`

<http:pgdfw-gopq/sql-bench-pgc-sh.txt> (<http:pgdfw-gopq/sql-bench-pgc-sh.txt>)

Apologies for txt links for source code

.txt file cause problems for browser - try hard refresh in browser

Software & Hardware Environment for Benchmark

PostgreSQL 9.4.1

4GB Shared Buffers
Compiled From Source

My Desktop iMac OSX 10.9.5

Intel Core i7 @ 3.4GHz, 4 Cores, 8 Threads
32G RAM, DDR3 @ 1600 MHz Bus, 8Mb L2 Cache
786Gb Apple SSD SM768E, Intel Controller

Query a Snapshot of Production HTTP Session Cookie Table

Tuple Width is 41 bytes

```
create table public.login_session
(
    cookie            udig primary key,    -- is just a sha1

    create_time       timestamp,
    active_time       timestamp
);
```

1,525,746 Tuples in Table with 160 bit Primary Key (SHA)

```
SELECT
    EXISTS (
        SELECT
            cookie
        FROM
            login_session
        WHERE
            cookie = $1
    )
;
```

Who Won the Shootout?

Go *pq* Package

16.50s for 1525861 queries = 92,476 lookup/sec

C *libpq* Library

15.37s for 1525861 queries = 99,275 lookup/sec

Perhaps pgbench in Go Language?

Links

[Interactive Tour of Go](http://tour.golang.org/) (http://tour.golang.org/)

[An Introduction to Programming in Go - Online Book](http://www.golang-book.com/) (http://www.golang-book.com/)

[Meet the Go Team - Q/A at Google I/O 2012](https://www.youtube.com/watch?v=sln-gJaURzk&list=PLoJWLKOp927tFsMbO2onhp26NnOAKAtO#t=921) (https://www.youtube.com/watch?v=sln-

gJaURzk&list=PLoJWLKOp927tFsMbO2onhp26NnOAKAtO#t=921)

[Concurrency is Not Parallelism by Rob Pike @ Vimeo](http://vimeo.com/49718712) (http://vimeo.com/49718712)

[Communication Sequential Processes - The Theory That Inspired Go @ Wikipedia](http://en.wikipedia.org/wiki/Communicating_sequential_processes)

(http://en.wikipedia.org/wiki/Communicating_sequential_processes)

[Communicating Sequential Processes - Readable Intro to CSP Algebra \(PDF\)](http://www.usingcsp.com/cspbook.pdf)

(http://www.usingcsp.com/cspbook.pdf)

[Searchable Documentation on Popular Go Packages](http://www.godoc.org) (http://www.godoc.org)

[PostgreSQL Package PQ Driver @ GoDoc](http://godoc.org/github.com/lib/pq) (http://godoc.org/github.com/lib/pq)

Thank you

April 14, 2015, Austin, Texas, USA

John Scott

Consultant, American Messaging, 2006-now

Founder, SetSpace, Inc, 1998-now

jmscott@setspace.com (mailto:jmscott@setspace.com)

john.scott@americanmessaging.net (mailto:john.scott@americanmessaging.net)

